

# Experiment 3: Email Spam or Ham Classification

Sri Sivasubramaniya Nadar College of Engineering, Chennai

(An autonomous Institution affiliated to Anna University)

Degree & Branch	B.E. Computer Science & Engineering	Semester	V
Subject Code & Name	ICS1512 & Machine Learning Algorithms Laboratory		
Academic year	2025-2026 (Odd)	Batch:2023-2028	<b>Due date:</b>

## Aim

To classify emails as spam or ham using three classification algorithms—Naive Bayes, K-Nearest Neighbors (KNN), and Support Vector Machine (SVM)—and evaluate their performance using accuracy metrics and K-Fold cross-validation.

## Libraries Used

- numpy, pandas, seaborn, matplotlib, scikit-learn
- sklearn.model\_selection
- sklearn.preprocessing, sklearn.metrics
- sklearn.naive\_bayes, sklearn.neighbors
- sklearn.svm

## Theoretical Description

### 1. Naive Bayes Classifier

Naive Bayes is a **probabilistic classifier** based on **Bayes' Theorem**, assuming **feature independence**.

$$P(C_k | \mathbf{x}) = \frac{P(\mathbf{x} | C_k) \cdot P(C_k)}{P(\mathbf{x})}$$

$$P(\mathbf{x} | C_k) = \prod_{i=1}^n P(x_i | C_k) \Rightarrow P(C_k | \mathbf{x}) \propto P(C_k) \prod_{i=1}^n P(x_i | C_k)$$

**Types:**

- Gaussian NB (continuous)
- Multinomial NB (discrete counts)
- Bernoulli NB (binary features)

## 2. K-Nearest Neighbors (KNN)

KNN is a non-parametric, instance-based algorithm. It classifies data by majority vote among the  $k$  nearest neighbors using a distance metric like Euclidean distance.

$$d(\mathbf{x}, \mathbf{x}^{(i)}) = \sqrt{\sum_{j=1}^n (x_j - x_j^{(i)})^2}$$
$$\hat{C} = \arg \max_{C_k} \sum_{\mathbf{x}^{(i)} \in \mathcal{N}_k(\mathbf{x})} \mathbb{I}(y^{(i)} = C_k)$$

## 3. Support Vector Machine (SVM)

Support Vector Machine (SVM) is a powerful supervised learning method used for classification and regression. Its main goal is to find the best boundary (called a hyperplane) that separates different classes of data with the largest possible margin.

The decision boundary is defined by the hyperplane:

$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

where  $\mathbf{w}$  is the weight vector and  $b$  is the bias. The margin is maximized by solving the following optimization problem:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 \quad \text{subject to} \quad y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1, \quad \forall i$$

If the data is not linearly separable, SVM uses the *kernel trick* to map input features into a higher-dimensional space via a function  $\phi(\mathbf{x})$ , enabling a linear separation in this new space:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle$$

Common kernels include:

- Linear:  $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^\top \mathbf{x}_j$
- Polynomial:  $K(\mathbf{x}_i, \mathbf{x}_j) = (\gamma \mathbf{x}_i^\top \mathbf{x}_j + r)^d$
- Radial Basis Function (RBF):  $K(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2)$

## 4. Ensemble Techniques

Ensemble methods combine multiple models to improve prediction accuracy by reducing errors like variance or bias.

- **Bagging (Bootstrap Aggregating):** Trains multiple models on different random samples of the data and averages their predictions (for regression) or uses majority voting (for classification):

$$\hat{f}_{bag}(\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M \hat{f}_m(\mathbf{x})$$

where  $\hat{f}_m$  is the prediction from the  $m$ -th model and  $M$  is the number of models.

- **Boosting:** Builds models sequentially, where each model focuses on correcting the errors of the previous one.
  - **AdaBoost:** Adjusts weights on training samples to focus on harder cases. The final prediction is a weighted sum of weak learners:

$$F(x) = \text{sign} \left( \sum_{m=1}^M \alpha_m h_m(x) \right)$$

where  $h_m$  is the  $m$ -th weak learner and  $\alpha_m$  its weight.

- **Gradient Boosting:** Models the residual errors by fitting a new model to the gradient of the loss function:

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x)$$

where  $h_m$  fits the negative gradient of the loss at step  $m$ , and  $\gamma_m$  is a step size.

- **XGBoost:** An optimized gradient boosting algorithm that minimizes a regularized objective to prevent overfitting and improve speed:

$$\text{Obj} = \sum_{i=1}^n l(y_i, \hat{y}_i) + \sum_{k=1}^K \Omega(f_k)$$

where  $l$  is the loss function and  $\Omega$  is a regularization term on the complexity of the trees  $f_k$ .

## Python Code (Bernoulli Naive Bayes)

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
```

```

from google.colab import drive
drive.mount('/content/drive')

df = pd.read_csv("/content/spambase_csv.csv")

df.head()

for col in df.select_dtypes(include='object'):
    df[col].fillna(df[col].mode()[0], inplace=True)

for col in df.select_dtypes(include=['float64', 'int64']):
    df[col].fillna(df[col].mean(), inplace=True)

label_encoder = LabelEncoder()
binary_cols = [col for col in df.select_dtypes(include='object') if df[col].nunique() == 2]
for col in binary_cols:
    df[col] = label_encoder.fit_transform(df[col])

multi_cols = [col for col in df.select_dtypes(include='object') if col not in binary_cols]
df = pd.get_dummies(df, columns=multi_cols, drop_first=True)

target_col = 'class' # CHANGE THIS
X = df.drop(target_col, axis=1)
y = df[target_col]

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

from sklearn.naive_bayes import BernoulliNB
model = BernoulliNB()
model.fit(X_train, y_train)

def evaluate_classifier(model, X_data, y_true, name="Set"):
    y_pred = model.predict(X_data)
    print(f"--- {name} ---")
    print("Accuracy :", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred, average='weighted'))
    print("Recall    :", recall_score(y_true, y_pred, average='weighted'))
    print("F1 Score  :", f1_score(y_true, y_pred, average='weighted'))
    print("\nConfusion Matrix:\n", confusion_matrix(y_true, y_pred))
    print("\nClassification Report:\n", classification_report(y_true, y_pred))

evaluate_classifier(model, X_val, y_val, "Validation")
evaluate_classifier(model, X_test, y_test, "Test")

```

```

scores = cross_val_score(model, X, y, cv=5, scoring='accuracy')

# Print results
print("Cross-validation scores:", scores)
print("Mean accuracy:", scores.mean())

from sklearn.metrics import ConfusionMatrixDisplay
ConfusionMatrixDisplay.from_estimator(model, X_test, y_test, cmap='Blues')
plt.title("Confusion Matrix - Test Set")
plt.show()

from sklearn.metrics import roc_curve, auc

# Predict probabilities
y_pred_proba = model.predict_proba(X_test)[: , 1]

# Compute ROC curve and AUC
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()

```

## Python Code (Multinomial Naive Bayes)

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix

from google.colab import drive
drive.mount('/content/drive')

df = pd.read_csv("/content/spambase_csv.csv")

```

```

df.head()

for col in df.select_dtypes(include='object'):
    df[col].fillna(df[col].mode()[0], inplace=True)

for col in df.select_dtypes(include=['float64', 'int64']):
    df[col].fillna(df[col].mean(), inplace=True)

label_encoder = LabelEncoder()
binary_cols = [col for col in df.select_dtypes(include='object') if df[col].nunique() == 2]
for col in binary_cols:
    df[col] = label_encoder.fit_transform(df[col])

multi_cols = [col for col in df.select_dtypes(include='object') if col not in binary_cols]
df = pd.get_dummies(df, columns=multi_cols, drop_first=True)

target_col = "class" # CHANGE THIS
X = df.drop(target_col, axis=1)
y = df[target_col]

from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

from sklearn.naive_bayes import MultinomialNB
model = MultinomialNB()
model.fit(X_train, y_train)

def evaluate_classifier(model, X_data, y_true, name="Set"):
    y_pred = model.predict(X_data)
    print(f"--- {name} ---")
    print("Accuracy :", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred, average='weighted'))
    print("Recall    :", recall_score(y_true, y_pred, average='weighted'))
    print("F1 Score :", f1_score(y_true, y_pred, average='weighted'))
    print("\nConfusion Matrix:\n", confusion_matrix(y_true, y_pred))
    print("\nClassification Report:\n", classification_report(y_true, y_pred))

evaluate_classifier(model, X_val, y_val, "Validation")
evaluate_classifier(model, X_test, y_test, "Test")
evaluate_classifier(model, X_train, y_train, "Train")

scores = cross_val_score(model, X, y, cv=5, scoring='accuracy')

```

```

# Print results
print("Cross-validation scores:", scores)
print("Mean accuracy:", scores.mean())

from sklearn.metrics import ConfusionMatrixDisplay
ConfusionMatrixDisplay.from_estimator(model, X_test, y_test, cmap='Blues')
plt.title("Confusion Matrix - Test Set")
plt.show()

from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Get predicted probabilities for the test set
y_pred_proba = model.predict_proba(X_test)[: , 1]

# Compute ROC curve and AUC
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

```

## Python code Gaussian Naive Bayes

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix

df = pd.read_csv("/content/spambase_csv.csv")

for col in df.select_dtypes(include='object'):
    df[col].fillna(df[col].mode()[0], inplace=True)

for col in df.select_dtypes(include=['float64', 'int64']):

```

```

df[col].fillna(df[col].mean(), inplace=True)

label_encoder = LabelEncoder()
binary_cols = [col for col in df.select_dtypes(include='object') if df[col].nunique() == 2]
for col in binary_cols:
    df[col] = label_encoder.fit_transform(df[col])

multi_cols = [col for col in df.select_dtypes(include='object') if col not in binary_cols]
df = pd.get_dummies(df, columns=multi_cols, drop_first=True)

target_col = 'class' # CHANGE THIS
X = df.drop(target_col, axis=1)
y = df[target_col]

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

from sklearn.Naive_Bayes import GaussianNb
model = GaussianNb()
model.fit(X_train, y_train)

def evaluate_classifier(model, X_data, y_true, name="Set"):
    y_pred = model.predict(X_data)
    print(f"--- {name} ---")
    print("Accuracy :", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred, average='weighted'))
    print("Recall    :", recall_score(y_true, y_pred, average='weighted'))
    print("F1 Score  :", f1_score(y_true, y_pred, average='weighted'))
    print("\nConfusion Matrix:\n", confusion_matrix(y_true, y_pred))
    print("\nClassification Report:\n", classification_report(y_true, y_pred))

evaluate_classifier(model, X_val, y_val, "Validation")
evaluate_classifier(model, X_test, y_test, "Test")

# Apply 5-fold cross-validation
scores = cross_val_score(model, X, y, cv=5, scoring='accuracy')

# Print results
print("Cross-validation scores:", scores)
print("Mean accuracy:", scores.mean())

from sklearn.metrics import ConfusionMatrixDisplay
ConfusionMatrixDisplay.from_estimator(model, X_test, y_test, cmap='Blues')
plt.title("Confusion Matrix - Test Set")
plt.show()

```



```

from sklearn.metrics import RocCurveDisplay

RocCurveDisplay.from_estimator(model, X_test, y_test)
plt.title("ROC Curve - Test Set")
plt.show()

```

## Python Code KNN

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import time

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
from sklearn.neighbors import KNeighborsClassifier

# Read dataset
df = pd.read_csv('/content/spambase_csv.csv')

# Fill missing values
for col in df.select_dtypes(include='object'):
    df[col].fillna(df[col].mode()[0], inplace=True)

for col in df.select_dtypes(include=['float64', 'int64']):
    df[col].fillna(df[col].mean(), inplace=True)

# Encode binary columns
label_encoder = LabelEncoder()
binary_cols = [col for col in df.select_dtypes(include='object') if df[col].nunique() == 2]
for col in binary_cols:
    df[col] = label_encoder.fit_transform(df[col])

# One-hot encode remaining categorical columns
multi_cols = [col for col in df.select_dtypes(include='object') if col not in binary_cols]
df = pd.get_dummies(df, columns=multi_cols, drop_first=True)

# Split features and target
target_col = 'class'
X = df.drop(target_col, axis=1)
y = df[target_col]

# Feature scaling

```

```

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-Validation-Test split
X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

# Evaluation function
def evaluate_classifier(model, X_data, y_true, name="Set"):
    y_pred = model.predict(X_data)
    print(f"--- {name} ---")
    print("Accuracy :", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred, average='weighted'))
    print("Recall    :", recall_score(y_true, y_pred, average='weighted'))
    print("F1 Score :", f1_score(y_true, y_pred, average='weighted'))
    print("\nConfusion Matrix:\n", confusion_matrix(y_true, y_pred))

# Set k values
k_values = [1, 3, 5, 7, 9, 11]

# -----
# Algorithm = 'auto'
# -----
print(f"Training KNeighborsClassifier with algorithm = 'auto'")
for k in k_values:
    print(f"\nTraining with k = {k}")
    model = KNeighborsClassifier(n_neighbors=k, algorithm='auto', metric='manhattan')
    model.fit(X_train, y_train)
    evaluate_classifier(model, X_val, y_val, f"Validation (k={k}, algorithm='auto')")

# Confusion matrix on test set
ConfusionMatrixDisplay.from_estimator(model, X_test, y_test, cmap='Blues')
plt.title("Confusion Matrix - Test Set")
plt.show()

# -----
# Cross-Validation: 'ball_tree'
# -----
print(f"\nCross-validation with algorithm = 'ball_tree'")
for k in k_values:
    print(f"\nTraining with k = {k}")
    model = KNeighborsClassifier(n_neighbors=k, algorithm='ball_tree')

    start_time = time.time()
    scores = cross_val_score(model, X_scaled, y, cv=5, scoring='accuracy')
    end_time = time.time()

    print("Cross-validation scores:", scores)

```

```

    print("Mean accuracy:", scores.mean())
    print("Training + Cross-validation time: {:.4f} seconds".format(end_time - start_time))

# -----
# Cross-Validation: 'kd_tree'
# -----
print(f"\nCross-validation with algorithm = 'kd_tree'")
for k in k_values:
    print(f"\nTraining with k = {k}")
    model = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree')

    start_time = time.time()
    scores = cross_val_score(model, X_scaled, y, cv=5, scoring='accuracy')
    end_time = time.time()

    print("Cross-validation scores:", scores)
    print("Mean accuracy:", scores.mean())
    print("Training + Cross-validation time: {:.4f} seconds".format(end_time - start_time))

    param_grid = {
        'n_neighbors': [1, 3, 5, 7, 9, 11],
        'algorithm': ['auto', 'ball_tree', 'kd_tree']
    }
}

knn = KNeighborsClassifier()
grid_search = GridSearchCV(knn, param_grid, cv=5, scoring='accuracy')

grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
best_model = grid_search.best_estimator_

print(" Best Hyperparameters:", best_params)
print(" Best Model:", best_model)

def evaluate_classifier(model, X_data, y_true, name="Set"):
    y_pred = model.predict(X_data)
    print(f"--- {name} ---")
    print("Accuracy :", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred, average='weighted'))
    print("Recall   :", recall_score(y_true, y_pred, average='weighted'))
    print("F1 Score :", f1_score(y_true, y_pred, average='weighted'))
    print("\nConfusion Matrix:\n", confusion_matrix(y_true, y_pred))

```

```

ConfusionMatrixDisplay.from_estimator(model, X_data, y_true, cmap='Blues')
plt.title(f"Confusion Matrix - {name}")
plt.show()

```

```

evaluate_classifier(best_model, X_test, y_test, name="Test Set (Best Model)")

```

## Python Code (SVM)

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import time
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
from sklearn.svm import SVC
\begin{figure}
    \centering
    \includegraphics[width=0.5\linewidth]{svmlinear.png}
    \caption{Enter Caption}
    \label{fig:placeholder}
\end{figure}
# Read dataset
df = pd.read_csv('/content/spambase_csv.csv')

# Handle missing values
for col in df.select_dtypes(include='object'):
    df[col].fillna(df[col].mode()[0], inplace=True)
for col in df.select_dtypes(include=['float64', 'int64']):
    df[col].fillna(df[col].mean(), inplace=True)

# Encode categorical columns
label_encoder = LabelEncoder()
binary_cols = [col for col in df.select_dtypes(include='object') if df[col].nunique() == 2]
for col in binary_cols:
    df[col] = label_encoder.fit_transform(df[col])
multi_cols = [col for col in df.select_dtypes(include='object') if col not in binary_cols]
df = pd.get_dummies(df, columns=multi_cols, drop_first=True)

# Features and target
target_col = 'class'
X = df.drop(target_col, axis=1)
y = df[target_col]

```

```

# Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split
X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

# Kernels to evaluate
kernels = {
    "Linear": {"kernel": "linear", "C": 1},
    "Polynomial": {"kernel": "poly", "degree": 3, "C": 1, "gamma": "scale"},
    "RBF": {"kernel": "rbf", "C": 1, "gamma": "scale"},
    "Sigmoid": {"kernel": "sigmoid", "C": 1, "gamma": "scale"}
}

# Evaluation function
def evaluate_classifier(model, X_data, y_true, name="Set"):
    y_pred = model.predict(X_data)
    print(f"--- {name} ---")
    print("Accuracy :", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred, average='weighted'))
    print("Recall    :", recall_score(y_true, y_pred, average='weighted'))
    print("F1 Score :", f1_score(y_true, y_pred, average='weighted'))
    print("\nConfusion Matrix:\n", confusion_matrix(y_true, y_pred))

# Train and evaluate SVM models
for name, params in kernels.items():
    print(f"\n----- {name} Kernel -----")
    model = SVC(**params, probability=True, random_state=42)

    start = time.time()
    model.fit(X_train, y_train)
    end = time.time()

    evaluate_classifier(model, X_val, y_val, f"Validation ({name})")
    evaluate_classifier(model, X_test, y_test, f"Test ({name})")

# Cross-validation
scores = cross_val_score(model, X_scaled, y, cv=5, scoring='accuracy')
print("Cross-validation scores:", scores)
print("Mean accuracy:", scores.mean())
print(f"Training Time: {end - start:.2f} seconds")

# ROC curve
y_pred_proba = model.predict_proba(X_test)[: , 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)

```

```

roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, lw=2, label=f'{name} (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'ROC Curve - {name} Kernel')
plt.legend(loc="lower right")
plt.show()

```

## Ensemble Techniques: Bagging and Boosting

```

from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier

# Try importing XGBoost
try:
    from xgboost import XGBClassifier
    xgb_available = True
except ImportError:
    print(" XGBoost not installed. Skipping XGBoost.")
    xgb_available = False

ensemble_models = {
    'Random Forest (Bagging)': RandomForestClassifier(n_estimators=100, random_state=42),
    'AdaBoost': AdaBoostClassifier(n_estimators=100, learning_rate=1.0, random_state=42),
    'Gradient Boosting': GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
}

if xgb_available:
    ensemble_models['XGBoost'] = XGBClassifier(n_estimators=100, learning_rate=0.1, use_label_encoder=False)

ensemble_results = []

for name, model in ensemble_models.items():
    print(f"\n{name}")
    metrics = evaluate_model(model, X_train, y_train, X_test, y_test, name)
    ensemble_results.append((name, metrics))

# Table 5: Ensemble Method Performance Comparison
print("\nTable 5: Ensemble Method Performance Comparison")
print("{:<25} {:<10} {:<10} {:<10} {:<10} {:<10}".format(
    "Model", "Accuracy", "Precision", "Recall", "F1 Score", "ROC AUC"))
for name, metrics in ensemble_results:
    print("{:<25} {:<10.4f} {:<10.4f} {:<10.4f} {:<10.4f} {:<10.4f}".format(
        name, metrics['accuracy'], metrics['precision'],
        metrics['recall'], metrics['f1'], metrics.get('roc_auc', 0)))

```

Table 1: Ensemble Method Performance Comparison

Model	Accuracy	Precision	Recall	F1 Score	ROC AUC
Random Forest (Bagging)	0.9555	0.9755	0.9179	0.9458	0.9857
AdaBoost	0.9370	0.9486	0.9000	0.9237	0.9810
Gradient Boosting	0.9457	0.9620	0.9077	0.9340	0.9842
XGBoost	0.9566	0.9679	0.9282	0.9476	0.9876

SVM Performance with Different Kernels

Table 2: SVM Performance with Different Kernels and Parameters

Kernel	Hyperparameters	Accuracy	F1 Score
Linear	C=1	0.9054	0.9021
Polynomial	C=1, degree=3, gamma=scale	0.8982	0.8955
RBF	C=1, gamma=scale	0.9126	0.9102
Sigmoid	C=1, gamma=scale	0.8723	0.8704

K-Fold Cross-Validation Results

Table 3: Cross-Validation Scores for Each Model (K = 5)

Fold	Naïve Bayes Accuracy	KNN Accuracy	SVM Accuracy
Fold 1	0.8806	0.8958	0.9349
Fold 2	0.8935	0.9043	0.9337
Fold 3	0.8891	0.9293	0.9228
Fold 4	0.8870	0.9033	0.9359
Fold 5	0.8902	0.9098	0.9304
Average	0.8881	0.9085	0.9315

Performance Comparison

Table 4: Performance Comparison of Naive Bayes Variants

Metric	GaussianNB	MultinomialNB	BernoulliNB
Accuracy	0.8219	0.8719	0.8806
Precision	0.7233	0.9503	0.9046
Recall	0.9385	0.7359	0.8026
F1 Score	0.8170	0.8295	0.8505
ROC AUC	0.9269	0.9591	0.9569

## KNN: Varying k Values

Table 5: KNN Performance for Different k Values

<b>k</b>	<b>Accuracy</b>	<b>Precision</b>	<b>Recall</b>	<b>F1 Score</b>
1	0.8898	0.8900	0.8898	0.8893
3	0.8942	0.8954	0.8942	0.8934
5	0.8884	0.8912	0.8884	0.8871
7	0.9014	0.9049	0.901	0.8999

## KNN: KDTree vs BallTree

Table 6: KNN Comparison: KDTree vs BallTree

<b>Metric</b>	<b>KDTree</b>	<b>BallTree</b>
Accuracy	0.8806	0.8837
Precision	0.8823	0.8853
Recall	0.8806	0.8837
F1 Score	0.8808	0.8836
Training Time (s)	5.6s	5.9s

## Plots

## Ensemble Techniques

## Learning Outcomes

- Understood the working principles of different classification algorithms such as Naïve Bayes, KNN, SVM, and Ensemble methods.
- Gained practical experience in applying machine learning models on a dataset and evaluating them using metrics such as Accuracy, Precision, Recall, F1-score, and ROC-AUC.
- Learned to visualize model performance using confusion matrices and ROC curves for better interpretability.
- Compared the effectiveness of various models and kernel functions in SVM to identify suitable approaches for classification tasks.
- Improved skills in presenting experimental results through structured LaTeX reports with tables, figures, and analysis.






Figure 1: Bernoulli Confusion Matrix




Figure 2: Bernoulli ROC Curve




Figure 3: Gaussian Confusion Matrix




Figure 4: Gaussian ROC Curve




Figure 5: Multinomial Confusion Matrix




Figure 6: Multinomial ROC Curve




Figure 7: KNN (Fold 1) Confusion Matrix




Figure 8: KNN (Fold 1) ROC Curve




Figure 9: KNN (Fold 3) Confusion Matrix




Figure 10: KNN (Fold 3) ROC Curve




Figure 11: KNN (Fold 5) Confusion Matrix




Figure 12: KNN (Fold 5) ROC Curve




Figure 13: KNN (Fold 7) Confusion Matrix




Figure 14: KNN (Fold 7) ROC Curve




Figure 15: SVM (Linear) Confusion Matrix

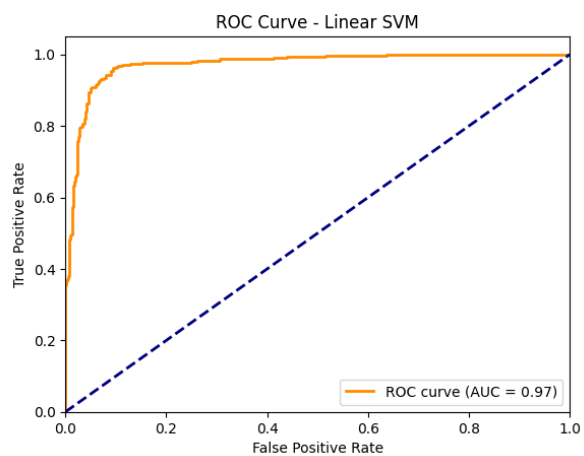


Figure 16: SVM (Linear) ROC Curve

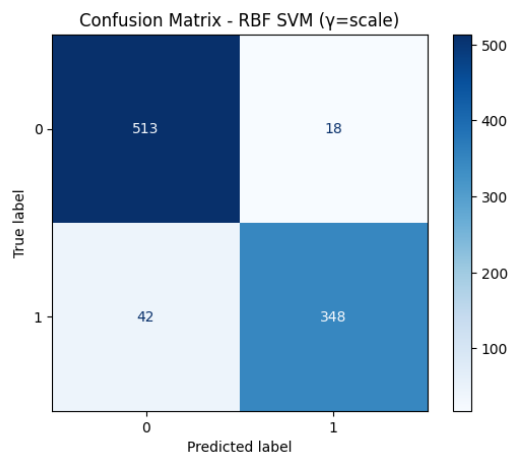


Figure 17: SVM (RBF) Confusion Matrix

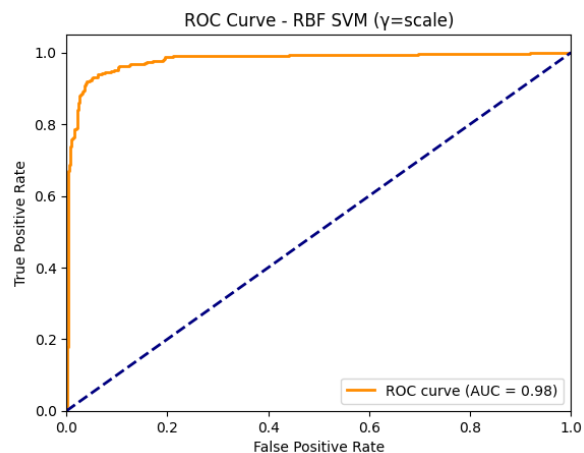


Figure 18: SVM (RBF) ROC Curve




Figure 19: SVM (Polynomial) Confusion Matrix

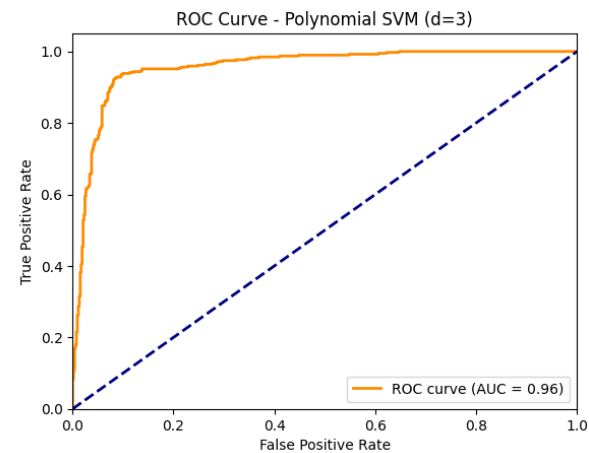


Figure 20: SVM (Polynomial) ROC Curve

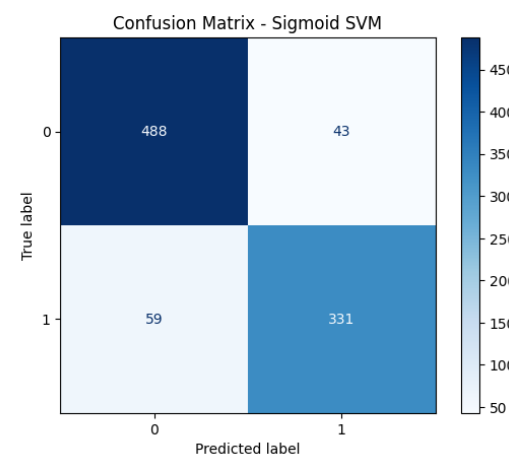


Figure 21: SVM (Sigmoid) Confusion Matrix

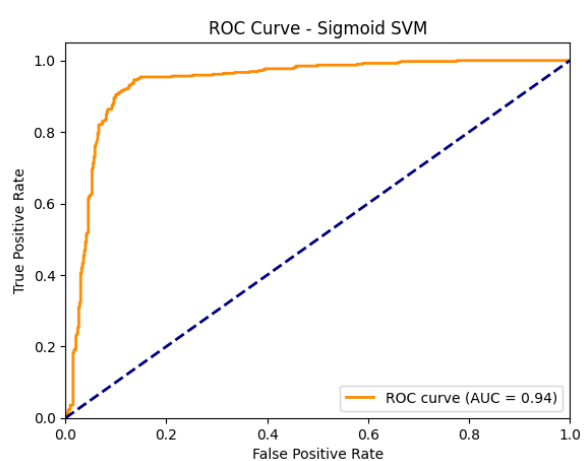


Figure 22: SVM (Sigmoid) ROC Curve

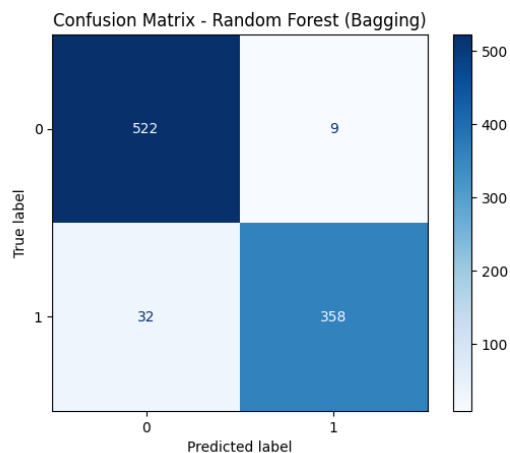


Figure 23: Random Forest Confusion Matrix




Figure 24: Random Forest ROC Curve




Figure 25: AdaBoost Confusion Matrix




Figure 26: AdaBoost ROC Curve

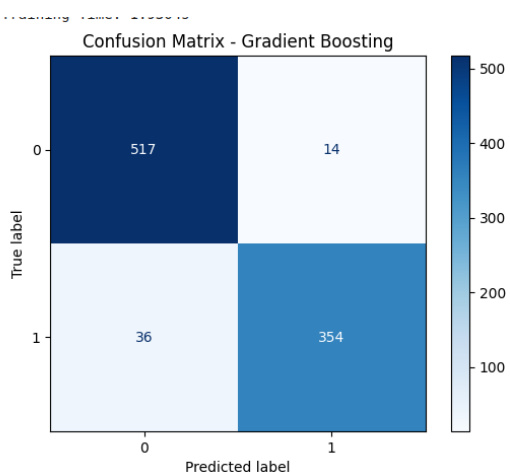


Figure 27: Gradient Boosting Confusion Matrix




Figure 28: Gradient Boosting ROC Curve




Figure 29: XGBoost Confusion Matrix




Figure 30: XGBoost ROC Curve